

Process Environment

Advanced Programming Environment in the UNIX Environment

Chun-Ying Huang <chuang@cs.nctu.edu.tw>

Outline

Process start and termination

Environment variables

Memory layout

Shared libraries

Memory allocation

setjmp and longjmp

Process resource limits

Process Start

The main function

Synopsis

- `int main(int argc, char *argv[]);`
- `int main(int argc, char *argv[], char *envp[]);`

Process Termination

Normal process termination in five ways

- Return from main
- Calling `exit`
- Calling `_exit` or `_Exit`
- Return of the last thread from its start routine
- Calling `pthread_exit` from the last thread

Abnormal process termination in three ways

- Calling `abort`
- Receipt of a signal
- Response of the last thread to a cancellation request

Execution of a main function looks like

- `exit(main(argc, argv));`

atexit and exit Functions

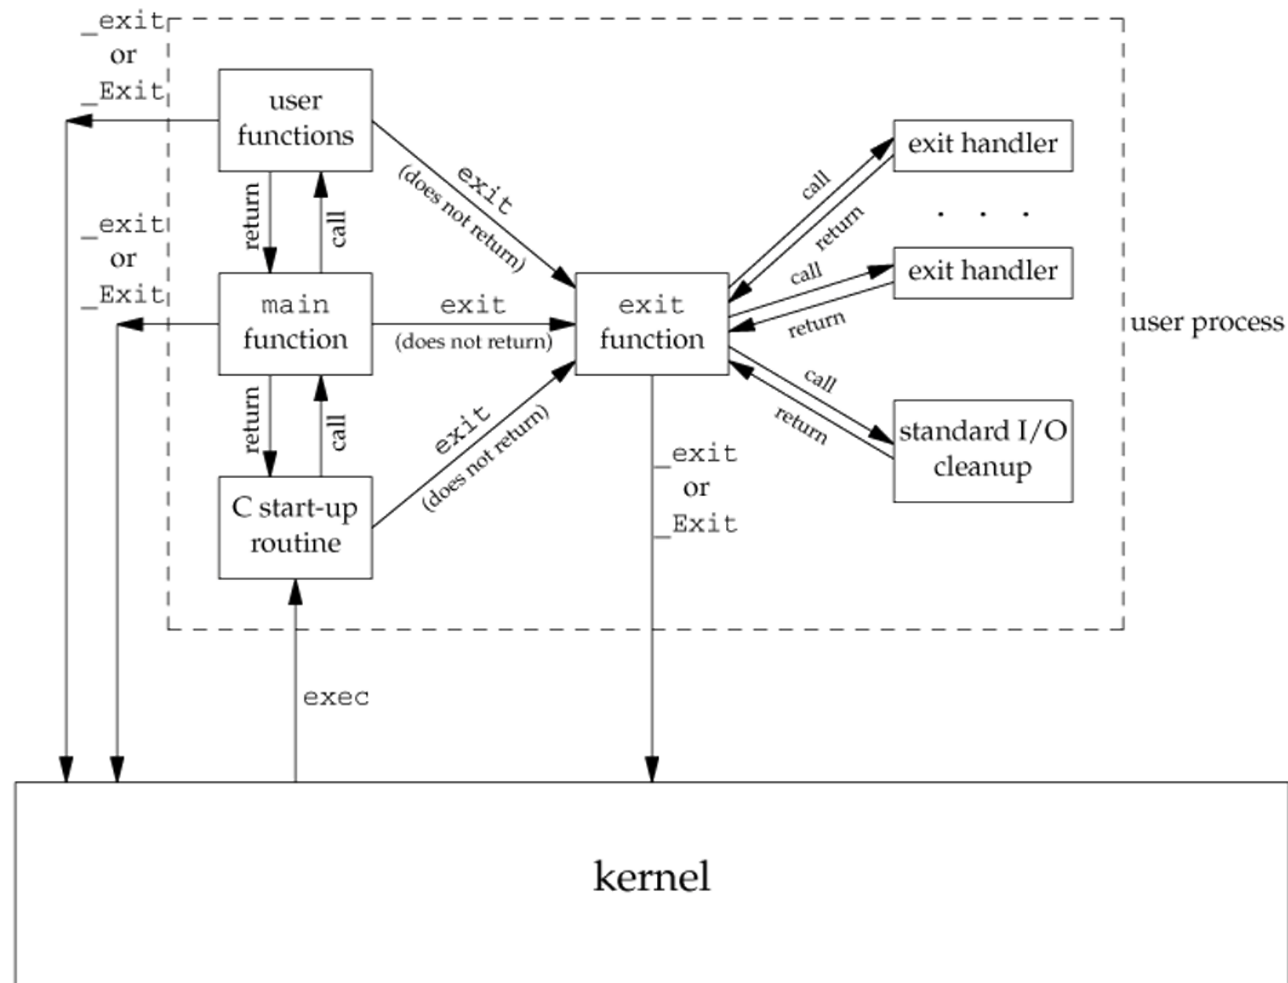
Manual cleanups on exit

- `int atexit(void (*function)(void));` *run before exit function*
- Register up to 32 customized functions (textbook) *也就是callback*
 - Linux has extended this restrictions

Exit functions

- `exit`
 - Call atexit registered functions
 - Performed a clean shutdown of the standard I/O library
 - `fclose()` all streams, `remove tmpfile()`
- `_exit` and `_Exit`
 - Terminate immediately

Start and Termination of a C Program



Environment Variables

The environment variables

- Usually in the form of: `name=value` (no spaces around =)
- Relevant commands: `env`, `export` (bash)
- Use `$` to read a specific environment variable in a shell

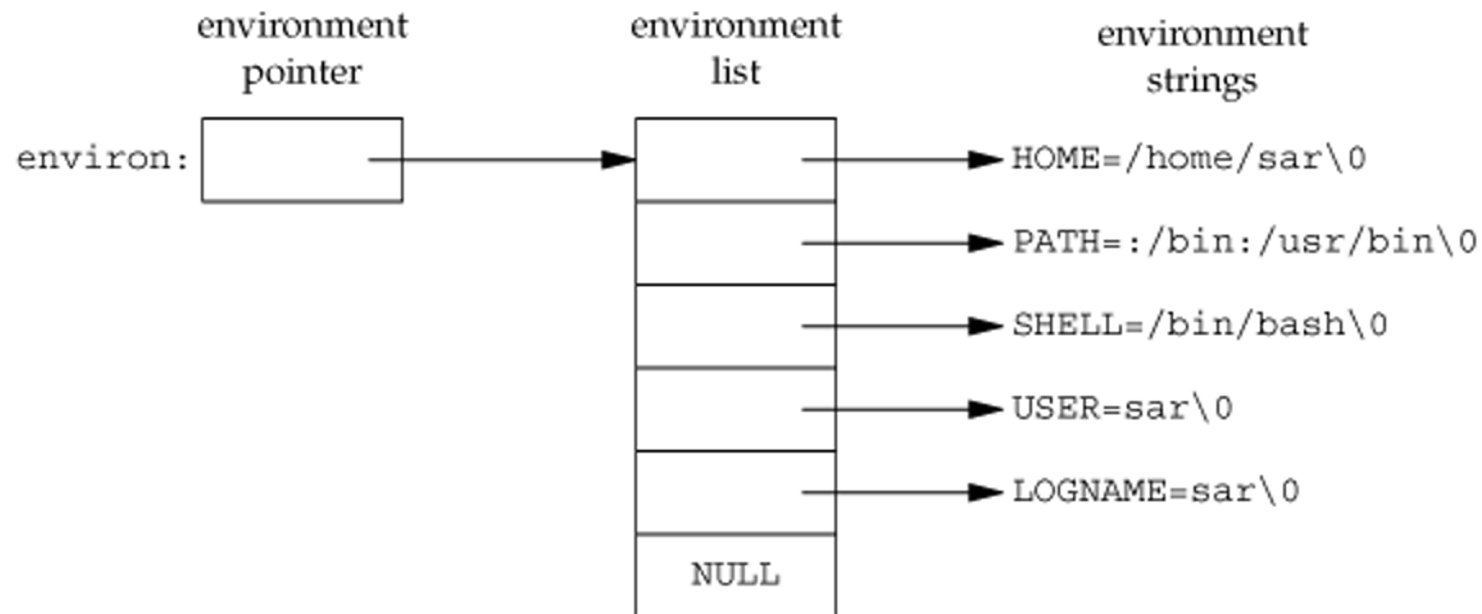
List of environment variable functions

| Function | ISO C | POSIX.1 | FreeBSD 8.0 | Linux 3.2.0 | Mac OS X 10.6.8 | Solaris 10 |
|-----------------------|-------|---------|----------------|----------------|--------------------|------------|
| <code>getenv</code> | • | • | • | • | • | • |
| <code>putenv</code> | | XSI | • | • | • | • |
| <code>setenv</code> | | • | • | • | • | |
| <code>unsetenv</code> | | • | • | • | • | |
| <code>clearenv</code> | | | | • | | |

Environment List

Access environment variables directly

- `int main(int argc, char *argv[], char *envp[]);`
- `extern char **environ;`



Environment Functions

Prototypes of functions to manipulate environment variables

```
#include <stdlib.h>
```

```
char *getenv(const char *name);  
int putenv(char *string);
```

```
int setenv(const char *name, const char *value, int overwrite);  
int unsetenv(const char *name);  
int clearenv(void);
```

Environment List Operations

Delete an entry

- This is simple, just free a string and move all subsequent pointers down one

Modify an entry

- If $\text{new-size} \geq \text{old-size}$, just overwrite the old one
- If $\text{new-size} > \text{old-size}$, allocate a new space the new variable and make the pointer point to the new location

Add an entry

- Add for the 1st time, allocate a new space for the entire list
- Add for non-1st time, reallocate a larger space for the entire list

Common Environment Variables (1/3)

| Variable | POSIX.1 | FreeBSD 8.0 | Linux 3.2.0 | Mac OS X 10.6.8 | Solaris 10 | Description |
|-------------|---------|-------------|-------------|-----------------|------------|---|
| COLUMNS | • | • | • | • | • | Terminal width |
| Datemask | XSI | | • | • | • | getdate(3) template file pathname |
| HOME | • | • | • | • | • | Home directory |
| LANG | • | • | • | • | • | Name of locale |
| LC_ALL | • | • | • | • | • | Name of locale |
| LC_COLLATE | • | • | • | • | • | Name of locale for collation |
| LC_CTYPE | • | • | • | • | • | Name of locale for character classification |
| LC_MESSAGES | • | • | • | • | • | Name of locale for messages |

Common Environment Variables (2/3)

| Variable | POSIX.1 | FreeBSD 8.0 | Linux 3.2.0 | Mac OS X 10.6.8 | Solaris 10 | Description |
|-------------|---------|-------------|-------------|-----------------|------------|--|
| LC_MONETARY | • | • | • | • | • | Name of locale for monetary editing |
| LC_NUMERIC | • | • | • | • | • | Name of locale for numeric editing |
| LC_TIME | • | • | • | • | • | Name of locale for date/time formatting |
| LINES | • | • | • | • | • | Terminal height |
| LOGNAME | • | • | • | • | • | Login name |
| MSGVERB | XSI | • | • | • | • | fmtmsg(3) message components to process |
| NLSPATH | • | • | • | • | • | Sequence of templates for message catalogs |

Common Environment Variables (3/3)

| Variable | POSIX.1 | FreeBSD 8.0 | Linux 3.2.0 | Mac OS X 10.6.8 | Solaris 10 | Description |
|----------|---------|-------------|-------------|-----------------|------------|---|
| PATH | • | • | • | • | • | List of path prefixes to search for executable file |
| PWD | • | • | • | • | • | Absolute pathname of current working directory |
| SHELL | • | • | • | • | • | Name of user's preferred shell |
| TERM | • | • | • | • | • | Terminal type |
| TMPDIR | • | • | • | • | • | Pathname of directory for creating temporary files |
| TZ | • | • | • | • | • | Time zone information |

Memory Layout of a Program

Text segment

- Machine instructions

Initialized data segment

- `int maxcount = 100;`

Uninitialized data segment (bss)

global/static `long sum[1000];`

Stack

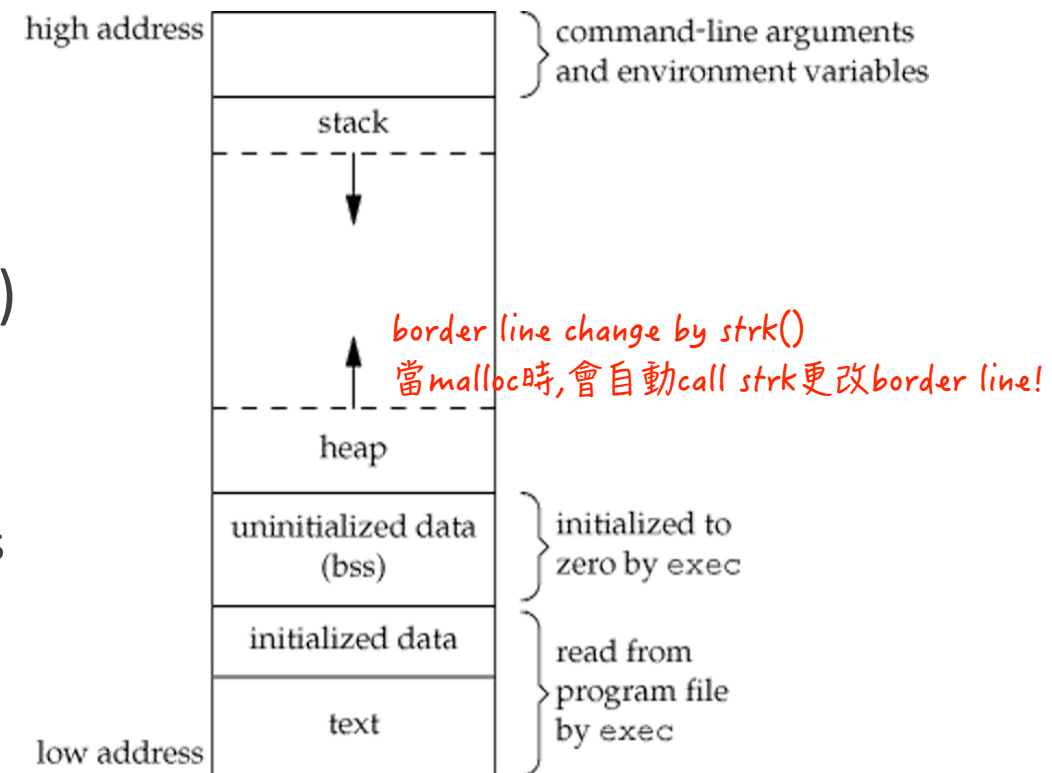
- Local variables, function call states

Heap

- Dynamic allocated memory

malloc

new



當 malloc 的 size 過大 (依據所定的 threshold) 時, 則不會把資料放在 heap, 而是放在其他地方

Read Sizes of an Executable Binary

The size (1) command

```
$ size /usr/bin/gcc /bin/sh
```

| text | data | bss | dec | hex | filename |
|--------|-------|-------|--------|-------|--------------|
| 203913 | 2152 | 2248 | 208313 | 32db9 | /usr/bin/gcc |
| 704028 | 19268 | 19736 | 743032 | b5678 | /bin/sh |

unix : .so windows : .dll mac : .dylib

Shared Libraries

Most UNIX systems today support shared libraries

Shared libraries remove the common library routines from the executable file

Maintain a **single copy of the library** routine somewhere in memory that all processes reference

- Reduce the size and memory requirement of each executable file
- But It may add some runtime overhead

Another advantage of shared libraries

- Library functions can be replaced with new versions without having to relink every program that uses the library
- But it might also be a security flaw

Compile Static and Dynamic Program

A simple program that just print “Hello, World!”

```
$ gcc h1.c -o h1
$ gcc h2.c -o h2 -static 不使用shared lib, 而是把所有用到的東西加載在text seg. 中
$ ls -la h1 h2
-rwxrwxr-x 1  chuang  chuang    9564 Mar 13 11:48 h1
-rwxrwxr-x 1  chuang  chuang  878192 Mar 13 11:48 h2
$ size h1 h2
   text    data     bss      dec     hex filename
   896     264        8    1168    490 h1
499650   1928    6948  508526  7c26e h2
```

Library Injection

Functions referenced to shared libraries can be overridden

- The LD_PRELOAD ^{environment variable} environment variable
- Usage:

LD_PRELOAD=/path/to/the/injected-shared-object {program}

在run {program} 之前, pre load自己所定義的lib, 查看是否有未解決的symbol在裡面

Library injection does not work with suid/sgid executables

nm [executable]: 查看該exec.的所有labels(symbols)以及對應的mem. 位址(似組語).

nm -D [executable]: 查看未解決的symbol, 即使用到dynamic lib的symbol.

file [executable]: 查看該exe. 是static linked or dynamic linked.

ldd [executable]: 查看該exe.的dy lib.

Library Injection Example

Suppose we are going to hijack the `getuid()` function

- This is commonly used in tools like fake-root

The original program (`getuid.c`)

```
int main() {  
    printf("UID = %d\n", getuid());  
    return 0;  
}
```

The injected library (`inject1.c`)

```
#include <stdio.h>  
#include <sys/types.h>  
  
uid_t getuid(void) {  
    fprintf(stderr, "injected getuid, always return 0\n");  
    return 0;  
}
```

Library Injection Example (Cont'd)

Compile the programs and the libraries

```
$ gcc -o getuid -Wall -g getuid.c
$ gcc -o inject1.so -shared -fPIC inject1.c -ldl
```

形成 shared lib

- The first command produces the getuid program
- The second commands generates the inject1.so (shared) library

Run the example

```
$ ./getuid                                # no injection
UID = 1000
$ LD_PRELOAD=./inject1.so ./getuid        # injected
injected getuid, always return 0
UID = 0
```

More on Library Injection

But we still want the original function to work properly

We have to locate the original function

當有自行撰寫的lib,但又想用原先系統提供的lib時

```
#include <dlfcn.h>
```

```
void *dlopen(const char *filename, int flag);  
char *dlerror(void);  
void *dlsym(void *handle, const char *symbol);  
int dlclose(void *handle);
```

You may have to link with `-ldl` option

dlopen()所用到的lib.(非傳統的clib,須額外開啟)

也就是在gcc injection.c時,需要加入這個參數 -ldl

Revised Library Injection Example

We would like to know the real UID internally (inject2.c)

```
#include <dlfcn.h>
#include <stdio.h>
#include <sys/types.h>

static uid_t (*old_getuid)(void) = NULL;    /* function pointer */

uid_t getuid(void) {
    if(old_getuid == NULL) {
        void *handle = dlopen("libc.so.6", RTLD_LAZY);
        if(handle != NULL)
            old_getuid = dlsym(handle, "getuid");
    }
    fprintf(stderr, "injected getuid, always return 0\n");
    if(old_getuid != NULL)
        fprintf(stderr, "real uid = %d\n", old_getuid());
    return 0;
}
```

Revised Library Injection Example (Cont'd)

Compile the programs and the libraries (again)

```
$ gcc -o getuid -Wall -g getuid.c
$ gcc -o inject2.so -shared -fPIC inject2.c -ldl
```

- The first command produces the getuid program
- The second commands generates the inject2.so (shared) library

Run the example

```
$ ./getuid                                # no injection
UID = 1000
$ LD_PRELOAD=./inject2.so ./getuid        # injected
injected getuid, always return 0
real uid = 1000
UID = 0
```

Determine Library Injection Possibility

No SUID/SGID enabled

Not a statically linked binary

Examples of the dynamic/static linked hello-world example

- The file command

```
$ file h1 h2
```

```
h1: ELF 64-bit LSB executable, x86-64, version 1 (SYSV), dynamically linked (uses shared  
libs), for GNU/Linux 2.6.24, BuildID[sha1]=e32f08cfbdda94d57273829c2bfd535d8fbe626d, not  
stripped
```

```
h2: ELF 64-bit LSB executable, x86-64, version 1 (GNU/Linux), statically linked, for  
GNU/Linux 2.6.24, BuildID[sha1]=2748d80822e76d183d0ef5633c0b784527727c7a, not stripped
```

- The ldd command

```
$ ldd h1 h2
```

```
h1:
```

```
linux-vdso.so.1 => (0x00007ffe7d3d5000)  
libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007f1bc2150000)  
/lib64/ld-linux-x86-64.so.2 (0x00007f1bc2515000)
```

```
h2:
```

```
not a dynamic executable
```


Determine Library Injection Possibility (Cont'd)

Use symbols from a shared library

The `nm` command

Example: static VS dynamic linked symbols

```
$ gcc -o getuid -Wall -g getuid.c                # dynamically linked
$ gcc -o getuid_s -Wall -g getuid.c -static      # statically linked
$ nm getuid | grep getuid
                 U getuid@@GLIBC_2.2.5          # getuid is unknown
$ nm getuid_s | grep getuid
000000000433590 W getuid                        # getuid is known (but weak)
000000000433590 T __getuid                      # the getuid implementation
```

Symbols can be stripped using the `strip` command

Memory Allocation

ISO C memory allocation functions

`void *malloc(size_t size);`

- Allocates a specified number of bytes of memory
- The initial value of the memory is indeterminate

`void *calloc(size_t nobj, size_t size);`

- Allocates space for a specified number of objects of a specified size
- The space is initialized to all 0 bits

`void *realloc(void *ptr, size_t newsize);`

- Increases or decreases the size of a previously allocated area
- It may involve moving the previously allocated area somewhere else, to provide the additional room at the end
- The initial value of increased memory is indeterminate

Memory Allocation (Cont'd)

Allocated memory can be released by `free()`

The allocation routines are usually implemented with the `sbrk(2)` system call
change the size of heap.

This system call expands (or contracts) the heap of the process

- However, most versions of `malloc` and `free` never decrease their memory size
- The space that we free is available for a later allocation
- The freed space is usually kept in the `malloc` pool, not returned to the kernel

The alloca Function

A special memory allocation function – alloca

```
#include <alloca.h>
void *alloca(size_t size);
```

alloca() allocate memories in **stack frames** of the current function call

So you don't have to free() the memory – it is released automatically after the execution of the current function returns

May be not supported by your system, but modern UNIXes supports the function (Linux, FreeBSD, Mac OS X, Solaris)

Pros: might be faster (than malloc), no need to free, easier to work with setjmp/longjmp
automatically restore when setjmp/longjmp

Cons: Portability

setjmp and longjmp Function

The reserved keyword "goto" can be used only in the same function

We cannot goto a label that is in another function

Instead, we must use the setjmp and longjmp functions to perform this type of branching

Typical Program Skeleton for Command Processing

```
#include "apue.h"

#define TOK_ADD    5

void do_line(char *);
void cmd_add(void);
int get_token(void);


int main(void) {
    char line[MAXLINE];
    while (fgets(line, MAXLINE, stdin) != NULL)
        do_line(line);
    exit(0);
}

char *tok_ptr; /* global pointer for get_token() */
void do_line(char *ptr) { /* process one line of input */
    int cmd;
    tok_ptr = ptr;
    while ((cmd = get_token()) > 0) {
        switch (cmd) { /* one case for each command */
            case TOK_ADD: cmd_add(); break;
        }
    }
}

void cmd_add(void) {
    int token;
    token = get_token(); /* rest of processing for this command */
}

int get_token(void) {
    /* fetch next token from line pointed to by tok_ptr */
}
```

What if we encounter an error in `cmd_add` and would like to jump back to the main function for processing the next line?



The Solution for Jumping Across Functions

Set the jump back position

- `int setjmp(jmp_buf env);`
global buffer: 存放當下 stack 狀態
- `env` is usually a global variable – has to be accessed from both the `setjmp` side and the `longjmp` side
- Returns: 0 if called directly, or nonzero if returning from a call to `longjmp`

Jump back

- `void longjmp(jmp_buf env, int val);`
 - The '`val`' will be returned from `setjmp`
 - If `val` is 0, it will be replaced by 1

Using setjmp and longjmp

```
#include "apue.h"
#include <setjmp.h>

#define TOK_ADD    5


jmp_buf jmpbuffer;

int main(void) {
    char line[MAXLINE];
    if (setjmp(jmpbuffer) != 0)
        printf("error");
    while (fgets(line, MAXLINE, stdin) != NULL)
        do_line(line);
    exit(0);
}

. . .

void cmd_add(void) {
    int token;
    token = get_token();
    if (token < 0) /* an error has occurred */
        longjmp(jmpbuffer, 1);
    /* rest of processing for this command */
}
```

Stack after jumped back



Restoration of Variables (1/4)

Type of variables

- Automatic, e.g., `[auto] int autoVal;`, the default
- Register, e.g., `register int regVal;`, store in register if possible
- Volatile, e.g., `volatile int volVal;`, store in memory

What are the values of variables after jumped back?

- It depends
- Most implementations do not try to roll back these automatic variables and register variables
- The standards say only that their values are indeterminate
- If you have an automatic variable that you do not want to be rolled back, define it with the volatile attribute
- Variables that are declared global or static are left alone when `longjmp` is executed
- **In short: variables in register – restored; variables in memory – kept**

Restoration of Variables (2/4)

```
#include "apue.h"
#include <setjmp.h>
static void    f1(int, int, int, int);
static void    f2(void);
static jmp_buf jmpbuffer;
static int     globval;

int main(void) {
    int         autoval;
    register int regival;
    volatile int volaval;
    static int   statval;
    globval = 1; autoval = 2; regival = 3; volaval = 4; statval = 5;
    if (setjmp(jmpbuffer) != 0) {
        printf("after longjmp:\n");
        printf("globval = %d, autoval = %d, regival = %d,"
               " volaval = %d, statval = %d\n",
               globval, autoval, regival, volaval, statval);
        exit(0);
    }
    // Change variables after setjmp, but before longjmp.
    globval = 95; autoval = 96; regival = 97; volaval = 98; statval = 99;
    f1(autoval, regival, volaval, statval); /* never returns */
    exit(0);
}

static void f1(int i, int j, int k, int l) {
    printf("in f1():\n");
    printf("globval = %d, autoval = %d, regival = %d,"
           " volaval = %d, statval = %d\n", globval, i, j, k, l);
    f2();
}

static void f2(void) { longjmp(jmpbuffer, 1); }
```

Restoration of Variables (3/4)

Rules for variable restoration

- Variables stored in memory will have values as of the time of calling `longjmp`
- Variables in the CPU and floating-point registers are restored to their values when `setjmp` was called

Hence,

- auto variables may be indeterminate, it depends on compiler implementations
- register variables are restored to the value of “before calling **setjmp**”
- volatile variable are restored to the value of “before calling **longjmp**”

Restoration of Variables (4/4)

Set 1,2,3,4,5 → setjmp → Set 95,96,97,98,99 → longjmp → ?

- No optimization: gcc places everything in memory
- Full optimization: auto/register variables are placed in registers

```
$ gcc fig7.13-testjmp.c -I../include -o t1
$ gcc fig7.13-testjmp.c -I../include -o t2 -O
$ ./t1
in f1():
globval = 95, autoval = 96, regival = 97, volaval = 98, statval = 99
after longjmp:
globval = 95, autoval = 96, regival = 97, volaval = 98, statval = 99
$ ./t2
in f1():
globval = 95, autoval = 96, regival = 97, volaval = 98, statval = 99
after longjmp:
globval = 95, autoval = 2, regival = 3, volaval = 98, statval = 99
```

local variable will store in register
任何variable都放在memory
compile without any optimization
compile with full optimization
特定variable會被放在register

Process Resource Limits

Every process has a set of resource limits

Resource limits are usually initialized by a parent process and inherited by its child processes

The `getrlimit` and `setrlimit` functions

```
#include <sys/time.h>
#include <sys/resource.h>
int getrlimit(int resource, struct rlimit *rlim);
int setrlimit(int resource, const struct rlimit *rlim);
```

The `rlimit` structure

```
struct rlimit {
    rlim_t rlim_cur; /* Soft limit */
    rlim_t rlim_max; /* Hard limit (ceiling for rlim_cur) */
};
```

Partial List of Process Resources

| Limit | XSI | FreeBSD 8.0 | Linux 3.2.0 | Mac OS X 10.6.8 | Solaris 10 |
|----------------|-----|-------------|-------------|-----------------|------------|
| RLIMIT_AS | • | • | • | | • |
| RLIMIT_CORE | • | • | • | • | • |
| RLIMIT_CPU | • | • | • | • | • |
| RLIMIT_DATA | • | • | • | • | • |
| RLIMIT_FSIZE | • | • | • | • | • |
| RLIMIT_MEMLOCK | | • | • | • | |
| RLIMIT_NOFILE | • | • | • | • | • |
| RLIMIT_NPROC | | • | • | • | |
| RLIMIT_RSS | | • | • | • | |
| RLIMIT_SBSIZE | | • | | | |
| RLIMIT_STACK | • | • | • | • | • |
| RLIMIT_VMEM | | | | | • |

Example to Dump Resource Limits

See code fig7.16-getrlimit.c

```
$ ./fig7.16-getrlimit
RLIMIT_AS          (infinite)  (infinite)
RLIMIT_CORE        1024000000 (infinite)
RLIMIT_CPU         (infinite)  (infinite)
RLIMIT_DATA        (infinite)  (infinite)
RLIMIT_FSIZE       (infinite)  (infinite)
RLIMIT_LOCKS       (infinite)  (infinite)
RLIMIT_MEMLOCK      65536      65536
RLIMIT_NOFILE       1024      4096
RLIMIT_NPROC        96120      96120
RLIMIT_RSS          (infinite)  (infinite)
RLIMIT_STACK        8388608  (infinite)
```

Example to Dump Resource Limits

| Limits | Description |
|----------------|--|
| RLIMIT_CORE | The maximum size in bytes of a core file. A limit of 0 prevents the creation of a core file. |
| RLIMIT_MEMLOCK | The maximum amount of memory in bytes that a process can lock into memory using mlock(2). |
| RLIMIT_NOFILE | The maximum number of open files per process. |
| RLIMIT_NPROC | The maximum number of child processes per real user ID. |
| RLIMIT_STACK | The maximum size in bytes of the stack. |

Q & A
